

NepPlural: A Human-Anchored Benchmark for Persona Annotation in Nepali Migration Discourse

Jinu Nyachhyon^{1,2}, Anish Khatiwada², Bimal Timilsina¹, Rakhee Pandey², Ashish Acharya^{2,3}

¹Institute for Research and Innovation in Intelligent Systems (IRIIS)

²Herald College Kathmandu (University of Wolverhampton)

³Islington College (London Metropolitan University)

Correspondence: nyachhyonjinu@gmail.com

Abstract

Pluralistic alignment asks that language models faithfully represent a society’s diverse, competing values rather than collapse them into one dominant viewpoint, yet flagship pluralistic-alignment datasets remain heavily skewed toward Western respondents. We introduce NepPlural, a human-anchored benchmark for modelling competing socio-cultural personas in Nepali public discourse on youth migration and brain drain. We curate 2,148 code-mixed (Devanagari, Romanized Nepali, English) YouTube comments, annotate them along a four-axis persona taxonomy (Intent, Primary Driver, Value Orientation, Affect) plus an explicit No-Persona class, and build a 300-comment gold set labelled independently by three native Nepali speakers (mean Fleiss’ $\kappa \approx 0.55$). Using this gold set to audit automatic annotation, we find that LLM annotation reliability is driven by protocol. Labels collected in batches through chat interfaces reach only near-chance inter-model agreement ($\kappa \approx 0.17$). Annotating one comment at a time raises agreement above the native-speaker level ($\kappa \approx 0.72$), but consistency is not validity: under a forced-choice schema, models agree on personas for the 39% of comments native speakers judge to express none, keeping macro-F1 against human gold near 0.50. Allowing abstention resolves this: models recover No-Persona comments with F1 of 0.79 to 0.87 and, against leave-one-annotator-out references, match or exceed individual native speakers (macro-F1 0.75 to 0.86 vs. 0.75). Encoders fine-tuned on the batched labels reach only 0.34. Finally, a human-judged steerability probe finds that three frontier models adopt single persona axes with 59 to 77% fidelity but compose full personas only 33% of the time, with no congruence gap detected at this scale between Western-congruent and Nepali-collectivist personas. We release the data, code, prompts, and all model outputs.¹

¹<https://github.com/iriis-research/NepPlural>

1 Introduction

As LLMs are deployed across the world’s linguistic communities, a central open problem is whether they can represent the diverse values of their users. Standard post-training alignment, such as RLHF (Bai et al., 2022) and DPO (Rafailov et al., 2023), optimizes toward a single aggregated response. This is appropriate for factual queries but corrosive for questions where reasonable people disagree (migration, family duty, national identity), producing “perspective collapse” toward a majority-Western annotator pool (Sorensen et al., 2024). Pluralistic alignment (Sorensen et al., 2024) asks instead that systems represent conflicting values rather than average them away. Work on cultural alignment has documented Western-centric value priors in LLMs (Cao et al., 2023; AlKhamissi et al., 2024; Naous et al., 2024), but largely through survey instruments, translated prompts, or curated probes rather than naturalistic discourse. Flagship pluralistic resources remain Western-skewed: PRISM includes almost no participants from South Asia and none from Nepal (Kirk et al., 2024), and GlobalOpinionQA (Durmus et al., 2024) and OpinionQA (Santurkar et al., 2023) derive from instruments designed in Western social science. For Nepali, existing NLP resources target lower-level tasks such as named-entity recognition and part-of-speech tagging (Subedi et al., 2024); resources for persona or value annotation do not exist.

We address this gap with NepPlural, a human-anchored benchmark for modelling competing socio-cultural personas in Nepali discourse on youth migration and brain drain. Migration is an apt testbed: it is among the most contested issues in contemporary Nepal, spanning economic necessity, family obligation, systemic anger, and patriotism, and it is debated online in exactly the code-mixed register (Devanagari, Romanized Nepali, English) that monolingual resources fail to cap-




Figure 1: Agreement is not validity. (a) Mean Fleiss’ κ among three native speakers and among three LLMs under each protocol. (b) Mean macro-F1 against leave-one-annotator-out references for a held-out native speaker and the LLM majority vote. Dashed lines mark the human level.

ture. Moving beyond binary sentiment, we annotate 2,148 comments along four axes (Intent, Primary Driver, Value Orientation, Affect) and an explicit No-Persona class for the off-topic reactions that saturate real comment sections.

Crucially, we do not treat LLM labels as ground truth. Annotation for low-resource languages increasingly relies on LLMs, with agreement across models used as a quality signal. We instead collect independent native-speaker labels for a 300-comment gold set and use them to audit automatic annotation under three protocols (Figure 1). Labels collected in batches through chat interfaces, a convenient practice when API access is limited, reach only near-chance inter-model agreement (Fleiss’ $\kappa \approx 0.17$, against 0.55 for native speakers). Annotating one comment at a time with the same instructions raises agreement to $\kappa \approx 0.72$, above the human level, yet the models then agree on personas for the 39% of comments native speakers judge to express none, and macro-F1 against human gold stays near 0.50; this pipeline would have passed a standard agreement check. Allowing the models to abstain resolves most of the gap: they recover No-Persona comments with F1 of 0.79 to 0.87, at a cost of 5 to 16% false abstentions, and against leave-one-annotator-out references they match or exceed individual native speakers (macro-F1 0.75 to 0.86 vs. 0.75).

Our contributions are:

- **The NepPlural dataset:** a code-mixed Nepali migration-discourse corpus with a four-axis persona taxonomy, a No-Persona class, and a 300-comment gold set from three native speakers.

- **A protocol audit of LLM annotation** that separates annotation format and access from abstention, showing that only human gold distinguishes agreement from validity.
- **An encoder benchmark** (five encoders, five seeds) that quantifies the downstream cost of an unreliable labelling protocol.
- **A human-judged steerability probe** showing that frontier models adopt single persona axes but rarely compose full personas.

We release the dataset, individual and adjudicated human labels, all annotation and generation prompts, raw model outputs for every condition, and the analysis code.

2 Related Work

2.1 Alignment and Steerable Pluralism

RLHF (Leike et al., 2018; Bai et al., 2022) and DPO (Rafailov et al., 2023) optimize toward a single aggregated notion of preference, which cannot capture values that conflict across individuals (Adams et al., 2025; Lindström et al., 2024). Sorensen et al. (2024) argue that this narrows the “Overton window” and distinguish Overton, steerable, and distributional pluralism. Steerable methods condition outputs on attribute profiles (Adams et al., 2025), combine community models (Feng et al., 2024), or personalize at inference time (Chen et al., 2025; Pădurean et al., 2025), while Sorensen et al. (2026) show that instruction tuning degrades in-context steerability, motivating tests in under-represented languages and cultures.

2.2 Cultural and Moral Alignment Evaluation

A growing literature, much of it gathered in the Cross-Cultural Considerations in NLP (C3NLP) workshop series, measures how well language models reflect the values of different societies. Cao et al. (2023) compare ChatGPT’s responses on Hofstede-style cultural survey items with human responses across countries, finding stronger alignment with American norms; Arora et al. (2023) probe pre-trained models for cross-cultural value differences across languages; and Jinnai (2024) asks whether cross-cultural alignment changes a model’s commonsense morality. AlKhamissi et al. (2024) show that prompting language and pretraining data composition shape how well models simulate survey responses from different countries, and Naous et al.

(2024) document Western-favouring cultural bias in models used for Arabic. Moral alignment likewise depends on the prompting language (Agarwal et al., 2024), motivating language-specific resources such as *Histoires Morales* for French (Leteno et al., 2025) and the MENA Values benchmark (Zahraei et al., 2025). Kharchenko et al. (2024) find that models can distinguish Hofstede’s dimensions across countries yet fail to apply them consistently in advice scenarios.

Table 1 positions NepPlural against representative resources. Most derive their items from survey instruments, often designed in Western social science, or from translated text; few use naturalistic discourse, and none offers human-gold persona labels with an explicit no-stance option in a low-resource, code-mixed language.

2.3 The Global-South Representation Gap

Even pluralistic resources inherit the Western skew of their contributor pools: PRISM’s South Asian coverage is thin and Nepal is absent (Kirk et al., 2024), and survey-derived benchmarks (Durmus et al., 2024; Santurkar et al., 2023) originate in Western social science. Gabriel et al. (2024) warn that aggregated supervision can encode culturally situated judgments as neutral principles, a concern acute for low-resource, culturally distinct languages. For Nepali, existing resources target tasks such as named-entity recognition and part-of-speech tagging (Subedi et al., 2024); to our knowledge, NepPlural is the first persona or value resource for the language (see Appendix A on Nepali honorific and pragmatic structure).

2.4 LLMs as Annotators

LLMs are increasingly used as annotators and judges (Gilardi et al., 2023; Zheng et al., 2023), motivated by cost and the scarcity of expert annotators for low-resource languages, with aggregation across samples or models as a standard safeguard (Wang et al., 2023). Prior work reports high LLM-human agreement on mainstream English tasks (Gilardi et al., 2023). Two assumptions underlie these safeguards: that agreement among models signals label quality, and that results are robust to how annotation is carried out. Our audit (§4) tests both for culturally situated labels in a low-resource language: inter-model agreement swings from near chance to above the human level with the annotation protocol alone, and high agreement can coexist with systematic over-attribution that only human

gold reveals. We therefore treat LLM labels as a signal to be measured against native-speaker annotation, not as ground truth.

3 Dataset and Taxonomy

3.1 Data Collection

To capture multi-perspective public discourse on migration, we collected comments from three high-engagement Nepali YouTube channels covering migration podcasts, interviews, and public commentary (IDS, Thaha Research, and The Nepali Comment), spanning a wide range of grassroots viewpoints. The result is informal, unedited stance-taking by ordinary participants rather than curated text.

3.2 Preprocessing and Anonymization

We removed author identifiers, timestamps, and platform comment IDs, replaced user mentions with a placeholder, and scrubbed other personal details (names, contact information), so that no comment is released in a form that identifies an individual. A regex-based routine normalizes Devanagari Unicode variants, and the IndicNLP Library segments sentences. A deliberately conservative filter drops a comment only if it carries no classifiable content (empty, emoji- or punctuation-only, fewer than two tokens, only URLs, mentions, or hashtags, spam, or exact duplicates). We apply no further lexical or script normalization, so Devanagari, Romanized Nepali, code-mixed, and English comments survive in their original form, each tagged with script, language, and source metadata. The final corpus contains 2,148 comments (Appendix B, Table 6).

3.3 Persona Taxonomy

We annotate each comment along four orthogonal axes that capture the persona and motivation behind a comment rather than its surface sentiment:

- **Intent** (stance): Pro-Migration, Anti-Migration, Trapped/Regretful, Neutral/Observation.
- **Primary Driver**: Economic Necessity, Family Obligation, Systemic/Political Anger, Patriotism/Love.
- **Value Orientation**: Collectivist-Family, Collectivist-Nation, Individualist-Self.

Resource	Language(s)	Item source	What is measured	Human gold	Naturalistic text
OpinionQA (Santurkar et al., 2023)	English	Pew public-opinion surveys (US)	Alignment with demographic opinion distributions	Survey	✗
GlobalOpinionQA (Durmus et al., 2024)	English	Pew Global Attitudes, World Values Survey	Alignment with country-level opinion distributions	Survey	✗
PRISM (Kirk et al., 2024)	Mostly English	Participant conversations with LLMs	Individual preferences and ratings	✓	Partial
Cao et al. (2023)	Multilingual	Hofstede-style survey items	Cultural dimension alignment	Survey	✗
AlKhamissi et al. (2024)	Arabic, English	World Values Survey items	Simulated survey responses by country	Survey	✗
Naous et al. (2024)	Arabic	Curated prompts and cultural entities	Western-favouring cultural bias	✓	Partial
HistoiresMorales (Leteno et al., 2025)	French	Translated, culturally adapted moral stories	Moral-norm alignment	✓	✗
NepPlural (ours)	Nepali (Devanagari, Romanized, code-mixed with English)	Public YouTube comments	Four-axis personas plus No-Persona; LLM annotation reliability; steerability	✓ (3 native speakers, adjudicated)	✓

Table 1: NepPlural compared with representative cultural, moral, and pluralistic alignment resources. “Survey” means ground truth comes from survey response distributions rather than item-level human labels.

- **Affect:** Despairing/Sad, Angry/Frustrated, Hopeful/Motivated, Pragmatic.

3.4 The No-Persona Class

Public comment sections are dominated by reactions directed at the video rather than the migration topic (generic praise, channel shout-outs, meta comments), for which no persona axis meaningfully applies. Rather than force labels onto such text, we introduce an explicit No-Persona class. A deterministic, auditable keyword-and-pattern router (no LLM) separates the corpus into a substantive pool and a No-Persona set: it routes a comment to No-Persona only on a positive match to praise or content-empty patterns, and never when a migration, political, economic, or family keyword is present; ambiguous cases stay in the substantive pool for human judgment (rules in Appendix C). This yields 1,894 substantive comments and 254 routed No-Persona comments. As §4 shows, native speakers judge a much larger share of the substantive pool to carry no persona as well, including comments that mention politics or money without taking any stance on migration (Table 2, rows e and f): under forced choice, all three LLMs assign personas to both; with abstention allowed, they correctly abstain on (e) but still over-attribute on (f).

4 Human Annotation and a Protocol Audit of LLM Annotation

Our central methodological commitment is that native-speaker judgment, not LLM consensus, defines ground truth for culturally situated Nepali personas. This section describes the human gold standard and then uses it to audit LLM annotation under three protocols.

4.1 Human Annotation Protocol

We draw a stratified sample of 300 comments from the substantive pool, near-proportional across source channels and preserving the script mix, with a light minimum quota so that rare personas are present in measurable numbers. Three native Nepali speakers independently annotated all 300 comments on the four axes following a shared guideline, and marked a comment No-Persona when it expressed no migration stance. Annotators did not confer on individual comments. Gold labels are the 2-of-3 majority per axis; the 72 cells with three-way ties were resolved by adjudication (class distribution in Appendix B, Table 7).

Because each annotator contributed to the gold labels, scoring an annotator against the gold overstates human performance (macro-F1 0.80). We therefore use a *leave-one-annotator-out* (LOO) reference: for each held-out annotator, the reference

	Comment (original)	English gloss	Human gold (I/D/V/A)	LLM labels by protocol
a	<i>rahar ley haina badhyata ley ho hajur</i>	It is not by choice but out of compulsion, sir.	Pro / Economic / Individualist / Despairing	–
b	<i>Manxe bides gaye ra xaena. desh xodi raxa hits hard</i>	It is not just that people went abroad; the country is being abandoned. It hits hard.	Anti / Patriotism / Coll.-Nation / Despairing	–
c	<i>Ani sarkar ko matra dosh chaina, sarkar ko under ma kam garne karmachari pani chorharule bhariyeko hunale pani ho</i> (orig. Devanagari)	Not only the government is to blame; its staff are full of thieves too.	Neutral / Systemic Anger / Coll.-Nation / Angry	–
d	<i>I want to go home but I can't</i>	(English)	Trapped / Family / Individualist / Despairing	–
e	<i>Euta video Lumbini Development Trust ko corruption ra politics ma ni banaunu paryo</i>	Someone should also make a video on corruption and politics at the Lumbini Development Trust.	No-Persona	Batched, forced: all three assign Neutral / Systemic Anger / Coll.-Nation. Abstain: all three No-Persona .
f	<i>Genuine saving kati samma hunx uta</i>	How much can one genuinely save over there?	No-Persona	Batched, forced: all three assign a persona (Economic / Individualist / Pragmatic in most). Abstain: all three still assign Pro / Economic / Individualist / Pragmatic.

Table 2: Examples from the gold set (emojis removed) across registers: Romanized Nepali (a), Romanized Nepali mixed with English (b), Devanagari with an English loanword (c), and English (d). Rows e and f are substantive-looking comments that native speakers judge to express no migration persona. I/D/V/A = Intent / Primary Driver / Value Orientation / Affect. Protocol labels are shown only for the No-Persona rows.

on an axis is the label on which the other two annotators agree, and items where they disagree are dropped for that axis (195 to 208 items per axis and rotation). Every system is scored on exactly the same items and references, and results are averaged over the three rotations. Under LOO, an individual native speaker reaches 0.75 mean macro-F1, which we use as the human reference throughout. LOO references cover the items on which humans converge, so comparisons against them concern consensus items rather than the most contested comments.

4.2 Human Agreement

Native speakers reach moderate agreement on every axis (mean Fleiss’ $\kappa = 0.55$; Table 3). Their agreement is highest on whether a comment carries a persona at all ($\kappa = 0.64$): native speakers judge 116 of the 300 comments (38.7%) to express no migration persona, even though the router kept all of them as substantive. Deciding *whether* a comment has a persona is easier than deciding *which*: on the 184 comments the gold marks as carrying a persona, agreement falls to $\kappa = 0.42$. We read this as genuine interpretive plurality in the discourse rather

than annotator noise, since the same comment can reasonably be heard as, for example, anger at the state or despair about it.

4.3 Three LLM Annotation Protocols

We compare three frontier models, GPT-5.2 (OpenAI, 2025), Gemini 3.5 Flash (Google DeepMind, 2026), and Claude Sonnet 4.6 (Anthropic, 2026), under three protocols, all using the same instructions and label definitions:

- **Batched chat:** our original pipeline. Comments were pasted into each model’s chat interface in numbered batches of 20 to 50, with a forced-choice instruction (“every comment gets all four tags”) and CSV output.
- **Per-comment, forced:** each comment is sent separately through an API (via OpenRouter, provider-default sampling), with the same forced-choice instruction and JSON output.
- **Per-comment, abstain:** identical to the previous protocol, except that the instruction lets the model mark a comment No-Persona instead of assigning labels.

	I	D	V	A	Mean	Pers.
Humans	0.56	0.58	0.55	0.52	0.55	0.42
Batched chat	0.17	0.13	0.18	0.22	0.17	0.24
Per-comment, forced	0.68	0.75	0.67	0.77	0.72	0.75
Per-comment, abstain	0.64	0.64	0.66	0.65	0.65	0.59

Table 3: Fleiss’ κ among three raters (three native speakers, or three LLMs under each protocol) per axis on all 300 gold comments, five-way label space including No-Persona. **Pers.:** mean κ on the 184 comments the gold marks as carrying a persona. On the No-Persona decision itself, κ is 0.64 for humans and 0.65 for the abstaining LLMs. I/D/V/A = Intent / Primary Driver / Value Orientation / Affect.

The last two differ only in the abstention option, so their contrast isolates abstention. The first two differ in several bundled respects (batching, chat interface versus API, CSV versus JSON, and default settings), so their contrast measures the combined effect of annotation format and access. The No-Persona definition given to the models adds two examples (generic praise, remarks addressed to the channel) beyond the human guideline wording. All prompts and raw outputs are released.

4.4 Agreement Depends on Protocol

Table 3 shows that inter-model agreement is a property of the protocol as much as of the models. Under batched chat annotation, the three models agree only slightly (mean $\kappa = 0.17$), and the gap to humans persists on persona-bearing comments (0.24 vs. 0.42), so it is not an artifact of forcing labels onto No-Persona text. Sending the same instructions one comment at a time raises agreement to $\kappa = 0.72$, above the native-speaker level, and to 0.75 on persona-bearing comments. Adding the abstention option leaves agreement similar (0.65), and the abstaining models agree with each other on the No-Persona decision about as well as humans do (0.65 vs. 0.64). Because the batched-to-per-comment contrast bundles several factors, we cannot attribute the gain to batching alone; what the audit does establish is that the near-chance agreement in our original pipeline reflected how annotation was carried out, not a ceiling on what these models can do.

4.5 Agreement Is Not Validity

High agreement among models does not mean their labels are right. Table 4 scores every model against human references. Under per-comment forced choice, the models agree with each other

		Gold	LOO	LOO-P	NP F1	FA (%)
	Human (held-out)	0.80 [†]	0.75	0.75	–	–
Batched	GPT	0.32	0.34	0.55	–	–
	Gemini	0.33	0.35	0.57	–	–
	Claude	0.39	0.40	0.67	–	–
	Majority	0.39	0.42	0.67	–	–
Forced	GPT	0.49	0.51	0.85	–	–
	Gemini	0.48	0.51	0.86	–	–
	Claude	0.47	0.49	0.83	–	–
	Majority	0.50	0.53	0.88	–	–
Abstain	GPT	0.70	0.80	0.79	0.83	15.8
	Gemini	0.75	0.86	0.86	0.87	4.9
	Claude	0.67	0.75	0.78	0.79	8.2
	Majority	0.74	0.85	0.86	0.88	5.4

Table 4: Validity against human references (mean macro-F1 over four axes). **Gold:** against the adjudicated gold, all 300 comments ([†]not independent for humans, who contributed to the gold). **LOO:** against leave-one-annotator-out references. **LOO-P:** LOO restricted to items where both reference annotators chose a persona. **NP F1:** No-Persona detection against the gold. **FA:** false-abstention rate on the 184 persona-bearing comments. Systems that cannot abstain score 0 on No-Persona items by construction; LOO-P removes that penalty.

more than native speakers do, yet reach only 0.47 to 0.50 mean macro-F1 against the gold, because they confidently assign personas to the 116 comments humans judge to express none. On items where both reference annotators chose a persona, the same labels score 0.83 to 0.88 (LOO-P): the persona judgments themselves are sound, and the failure is over-attribution. A pipeline that checked quality by inter-model agreement would have accepted these labels; only the human gold reveals the error.

Allowing abstention resolves most of it. The abstaining models recover the No-Persona comments with F1 of 0.79 to 0.87 against the gold, at a cost of 4.9% (Gemini) to 15.8% (GPT) false abstentions on persona-bearing comments. Against LOO references, they reach 0.75 to 0.86 mean macro-F1 (majority vote 0.85), matching or exceeding a held-out native speaker (0.75), and they stay at 0.78 to 0.86 on persona-only references, so the gain does not come at the expense of persona labels. Abstention does not remove over-attribution entirely: all three abstaining models still label a question about savings abroad as Pro-Migration (Table 2, row f).

4.6 Results by Script

Broken down by script (Appendix D, Table 8), the abstaining majority matches or exceeds a held-out

native speaker in every group, while batched annotation and the encoders are weakest on code-mixed comments. The Devanagari and code-mixed groups are small (11 and 8 LOO reference items on average), so this is indicative only. Separately, an earlier LLM-as-a-Judge check of a 20% sample of the batched labels flagged 35.2% of rows, concentrated in off-topic comments (Appendix E).

5 Encoder Benchmark

We measure how far fine-tuned encoders get when trained on cheap automatic labels and evaluated against human references.

5.1 Setup

The 300-comment gold set is the test set. Training and validation draw on the remainder of the corpus with no overlap: the batched-chat LLM majority labels (§4) on the substantive pool for the four persona axes, and the router’s No-Persona comments as the No-Persona signal, giving 1,661 training and 184 validation examples. These labels predate the audit, so the encoders inherit the batched protocol’s unreliability; this section therefore measures what that protocol yields downstream. We fine-tune five encoders as multi-task classifiers that jointly predict all four axes over a five-way label space per axis: multilingual BERT (Devlin et al., 2019), XLM-RoBERTa base (Conneau et al., 2020), two Nepali-specific models from IRIIS (BERT 110M, RoBERTa 125M; Thapa et al., 2025), and NepBERTa (Timilsina et al., 2022). Each is trained over five seeds (learning rate 10^{-5} , batch size 16, maximum length 192, 10 epochs, best-validation checkpoint); we report means over seeds.

5.2 Results

Table 5 reports macro-F1 against the gold and against the LOO references of §4. The multilingual encoders lead (0.34 against the gold, 0.35 to 0.36 LOO), and the Nepali-specific models trail at about 0.20. This gap is confounded with pretraining scale: the multilingual models saw far more data, so the difference should not be read as an effect of multilinguality itself.

All encoders fall far below a held-out native speaker (0.75 LOO), but the reference rows show why this gap should not be read as task difficulty alone. The best encoders score close to the labels they were trained on (batched majority: 0.39 against the gold), while the same three

	Gold	LOO	LOO-P
<i>References</i>			
Human (held-out)	0.80 [†]	0.75	0.75
Abstain LLMs, zero-shot (majority)	0.74	0.85	0.86
Batched LLMs (majority; training labels)	0.39	0.42	0.67
<i>Fine-tuned encoders (mean of 5 seeds)</i>			
XLM-RoBERTa base	0.34	0.36	0.49
mBERT (cased)	0.34	0.35	0.49
NepBERTa	0.21	0.20	0.24
IRIIS RoBERTa 125M	0.20	0.19	0.24
IRIIS BERT 110M	0.20	0.19	0.24

Table 5: Encoder results (mean macro-F1 over four axes) on the 300 gold comments, with the metrics of Table 4. Encoders are trained on batched LLM labels; the reference rows show the quality of those labels and of zero-shot LLMs under the abstention protocol. [†]Not independent for humans. The multilingual vs. Nepali-specific gap is confounded with pretraining scale.

LLMs, queried zero-shot under the abstention protocol, reach 0.74 against the gold and 0.85 LOO. The binding constraint is thus largely the quality of the training labels; retraining the encoders on abstention-protocol labels is a direct next step.

Per axis (Appendix F, Table 10), Value Orientation is the most learnable (up to 0.43) and Intent the least (0.27 to 0.28). No-Persona is the hardest class for every encoder (per-class F1 0.20 to 0.28): the router supplies only a narrow No-Persona training signal (overt praise and content-empty text), whereas the human-judged No-Persona test cases include substantive-looking comments that take no stance. We do not interpret per-class scores for classes with fewer than about 20 test examples.

6 Steerability Evaluation

We now use NepPlural for its motivating purpose: probing whether generative LLMs can be steered to adopt Nepali migration personas, and whether they do so evenhandedly. Steerable pluralism (Adams et al., 2025; Sorensen et al., 2024) requires that a model, given an in-context specification of a perspective, produces output that faithfully expresses it. Following reports that instruction tuning collapses models toward generic defaults (Sorensen et al., 2026), a natural worry is that frontier models steer more readily toward personas congruent with Western-centric defaults than toward Nepali-collectivist ones, a *congruence gap*.

6.1 Design

We define six target personas as four-axis bundles, each with a short natural-language description (Ap-

pendix G, Table 12). We call a persona **Western-congruent** when it is individualist (Individualist-Self), treats migration as a personal economic calculation (Economic Necessity), and is emotionally detached (Pragmatic), matching the individualist, mobility-positive framing that cultural-alignment studies associate with Western defaults (Cao et al., 2023; AlKhamissi et al., 2024). We call a persona **Nepali-collectivist** when its value orientation is collectivist (toward nation or family) and its stance is grounded in patriotism, family duty, or anger at the state. Congruence is thus carried by Value Orientation, our most reliably annotated axis, but it co-varies with Primary Driver and Affect, and two congruent personas share a target bundle; the comparison is between these two persona families, not between value orientations in isolation.

We pair the personas with six neutral migration questions and prompt three frontier models, Claude Opus 4.6, GPT 5.6 Luna, and Gemini 3.6 Flash, to answer as each persona in the code-mixed register of the corpus, using one prompt phrasing per persona and temperature 1.0 for every model. This yields 108 generations (3 models $\times$ 6 personas $\times$ 6 questions), 54 per persona family.

6.2 Human-Judged Fidelity

Because using an LLM judge would reintroduce the circularity this paper argues against, and because our encoders are only moderately accurate (§5), we score fidelity by human annotation. The same three native speakers labelled all 108 generations on the four axes, blind to target persona, model, and persona family (generations shuffled, targets hidden). Fidelity on an axis means the human-consensus label matches the target; strict fidelity requires all four axes to match.

6.3 Results

Two findings emerge. First, models steer individual axes moderately well, with fidelity of 59.3% on Intent, 75.0% on Primary Driver, 69.4% on Value Orientation, and 76.9% on Affect, but compose the full four-axis persona only 33.3% of the time. Adopting one specified attribute is achievable; holding four at once is not, which is the regime a pluralistic system must handle, since real personas are bundles of co-occurring values.

Second, we detect no congruence gap. Fidelity on Value Orientation is 68.5% for Western-congruent and 70.4% for Nepali-collectivist personas (gap -1.9 points, 95% bootstrap CI

$[-20.4, +16.7]$), and strict fidelity is 33.3% for both (CI $[-18.5, +18.5]$); permutation tests give $p = 1.0$. The confidence intervals are wide: with 54 generations per family, one phrasing per persona, and one temperature per model, we can rule out only gaps of roughly 20 points or more, and a smaller gap, or one that appears under other prompts or at larger scale, remains possible. Per-model differences (Appendix H) are inconsistent in sign, from $+16.7$ points (GPT) to -22.2 (Gemini) on 18 generations per cell, and we treat them as noise. We therefore report a bounded negative result: at this scale, no large directional bias against Nepali-collectivist personas is detectable, and the clearer limitation is compositional, holding several persona attributes at once.

7 Discussion

Annotation reliability is a property of the protocol. The same three models agree at near chance when annotating batches in chat interfaces and above the native-speaker level when annotating one comment at a time with the same instructions (§4). Low inter-model agreement therefore need not mean that models cannot perform a task, and our own earlier reading of it as such was wrong. Batching, interface, output format, and default settings belong in the methods section of any LLM-annotated resource, and a pipeline should be audited before its labels are trusted or used for training.

Agreement is not validity. Under per-comment forced choice, models agreed more than native speakers while assigning personas to every comment humans judged to carry none. A quality check based on inter-model agreement would have passed these labels. For culturally situated annotation, human gold is needed in both directions: to expose a pipeline that looks broken for reasons unrelated to the task, and to catch one that looks healthy but is systematically wrong. With a careful protocol, frontier models matched individual native speakers on consensus items, but only the human gold could establish that.

Abstention is a first-class capability. Allowing models to say that a comment expresses no persona turned consistent labels into valid ones, raising macro-F1 against the gold from about 0.50 to between 0.67 and 0.75. For pluralistic representation, knowing when not to speak for someone

matters as much as knowing which perspective to voice. Abstention is not free (5 to 16% false abstentions) and does not remove over-attribution entirely (Table 2, row f), so the abstention behaviour is itself a design choice that should be reported and validated.

Human disagreement carries signal. Native speakers agree strongly on whether a comment carries a persona but only moderately on which one ($\kappa = 0.42$ on persona-bearing comments), while per-comment LLMs agree more with each other (0.59 to 0.75). Higher model agreement may partly reflect shared priors rather than better judgment, and for a pluralistic resource the spread of native-speaker readings is informative in its own right. We release the individual annotations, not only the adjudicated gold, to support distributional rather than single-label evaluation.

8 Conclusion

We introduced NepPlural, a human-anchored benchmark for competing socio-cultural personas in Nepali migration discourse, and used its gold standard to audit LLM annotation. LLM annotation reliability turned out to depend on protocol: batched chat annotation yields near-chance agreement, per-comment annotation yields high agreement that hides over-attribution, and per-comment annotation with an abstention option reaches native-speaker level on consensus items. In persona steering, frontier models adopt single attributes moderately well but rarely compose full personas, with no large bias against Nepali-collectivist personas detected at our scale. We release the dataset, individual and adjudicated human labels, all prompts and raw model outputs, the encoder baselines, and the steerability set. Future work includes retraining encoders on abstention-protocol labels, adding open-weight models, enlarging the annotator pool, extending to other contested Nepali domains, and scaling the steerability study across prompt phrasings and sampling settings.

Limitations

Scale and domain. The corpus is single-domain (migration and brain drain), single-platform (YouTube), and drawn from three channels, so findings may not transfer to other topics or platforms. The gold set has 300 comments, several classes have fewer than 20 examples, and the Devanagari and code-mixed script groups are small; per-class

and per-script results are indicative only. Script tags come from a heuristic detector that misses some Romanized Nepali.

Annotators and the human reference. Three annotators suffice for majority-vote gold and agreement statistics, but they were members of the research effort rather than independently recruited, and a larger, demographically broader pool would strengthen the ground truth. Because annotators contributed to the gold, we report a leave-one-annotator-out reference (0.75) rather than the non-independent 0.80. LOO references cover only items on which two annotators agree, so comparisons against them concern consensus items, not the most contested comments.

Protocol comparison. The batched-chat and per-comment protocols differ in several bundled respects (batching, chat interface versus API, CSV versus JSON output, and default settings), so we cannot attribute their difference to any single factor. Only the forced-versus-abstain contrast is a controlled comparison. The No-Persona definition given to the models includes two examples beyond the human guideline wording, and the per-comment runs used provider-default sampling through OpenRouter.

Model coverage. All annotators and generators are proprietary models. We did not evaluate open-weight generative LLMs, so we cannot say whether the protocol effects or the steerability results extend to them; adding multilingual open-weight models is a priority for reproducible comparison. Proprietary models are updated over time, so exact reproduction depends on versions we record but do not control; the dataset, human labels, and released model outputs are the fixed artifacts.

Model limitations versus training-label noise. The encoders were trained on the batched LLM labels, which the audit shows to be unreliable, and evaluated on human gold. Their scores therefore mix three factors: noise in the training labels, the mismatch between those labels and human judgment, and model capacity. Zero-shot LLMs under the abstention protocol score far higher on the same test set, which suggests that label quality is the main constraint, but separating the factors properly requires encoders trained on human or abstention-protocol labels, which we leave to future work. The router’s narrow No-Persona training signal further depresses that class.

Steerability scale. The steerability study uses 108 generations (54 per persona family), one prompt phrasing per persona, and one temperature per model, with different model versions from the annotation study. Its confidence intervals rule out only gaps of roughly 20 points or more. Persona family is confounded with Primary Driver and Affect targets, and two congruent personas share a target, so the result is a bounded negative finding about these persona families, not about value orientation in isolation.

Ethics Statement

Data and privacy. NepPlural is derived from public YouTube comments. We removed author identifiers, timestamps, and platform comment IDs, replaced user mentions with a placeholder, and scrubbed other personal details; only anonymized text is released. Verbatim public comments can in principle be located by search even without author handles, and we ask users of the dataset not to attempt re-identification. Comment text was sent to commercial model providers for annotation and generation under their API and interface terms.

Sensitivity. Migration and brain drain are politically charged in Nepal. Our labels describe the persona a comment expresses, not the truth of any claim, and must not be used to profile, target, or make consequential judgments about individuals or communities. A data statement (Bender and Friedman, 2018) accompanies the release.

Annotation. The three annotators are native Nepali speakers familiar with the discourse under study. They annotated as part of the research effort rather than as independently recruited or compensated participants. Annotation guidelines are released with the dataset.

Intended use. NepPlural is intended for research on pluralistic alignment, low-resource NLP, LLM annotation methodology, and cultural representation in language models. It is not suitable for deployment decisions about individuals.

Acknowledgments

We thank the anonymous reviewers for detailed and constructive feedback; in particular, the request for an abstention-enabled baseline led directly to the protocol audit at the center of this version.

Declaration on Generative AI Use

Generative AI models are a direct object of study in this work: frontier LLMs serve as annotators under three protocols (§4), a separate model serves as an LLM-as-a-Judge verifier, and three models generate persona-conditioned text (§6). All such uses are reported with model identifiers and settings. Generative AI tools were also used to assist with structuring the manuscript and drafting analysis scripts. All scripts were reviewed and their outputs checked against the underlying data, and all claims, results, and analysis are the responsibility of the human authors.

References

- Jadie Adams, Brian Hu, Erik Veenhuis, David Joy, Bharadwaj Ravichandran, Aaron Bray, Anthony Hoogs, and Arslan Basharat. 2025. [Steerable pluralism: Pluralistic alignment via few-shot comparative regression](#). In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society (AIES)*, volume 8, pages 15–25. AAAI Press.
- Utkarsh Agarwal, Kumar Tanmay, Aditi Khandelwal, and Monojit Choudhury. 2024. Ethical reasoning and moral value alignment of LLMs depend on the language we prompt them in. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*.
- Badr AlKhamissi, Muhammad ElNokrashy, Mai AlKhamissi, and Mona Diab. 2024. Investigating cultural alignment of large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics.
- Anthropic. 2026. [System card: Claude Sonnet 4.6](#). Technical report, Anthropic.
- Arnav Arora, Lucie-aimée Kaffee, and Isabelle Augenstein. 2023. Probing pre-trained language models for cross-cultural differences in values. In *Proceedings of the First Workshop on Cross-Cultural Considerations in NLP (C3NLP)*. Association for Computational Linguistics.
- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, Nicholas Joseph, Saurav Kadavath, Jackson Kernion, Tom Conerly, Sheer El-Showk, Nelson Elhage, Zac Hatfield-Dodds, Danny Hernandez, Tristan Hume, and 12 others. 2022. [Training a helpful and harmless assistant with reinforcement learning from human feedback](#). Preprint, arXiv:2204.05862.
- Emily M. Bender and Batya Friedman. 2018. [Data statements for natural language processing: Toward](#)

- mitigating system bias and enabling better science. *Transactions of the Association for Computational Linguistics (TACL)*, 6:587–604.
- Yong Cao, Li Zhou, Seolhwa Lee, Laura Cabello, Min Chen, and Daniel Hershcovich. 2023. Assessing cross-cultural alignment between ChatGPT and human societies: An empirical study. In *Proceedings of the First Workshop on Cross-Cultural Considerations in NLP (C3NLP)*. Association for Computational Linguistics.
- Ruizhe Chen, Xiaotian Zhang, Meng Luo, Wenhao Chai, and Zuozhu Liu. 2025. PAD: Personalized alignment of LLMs at decoding-time. In *International Conference on Learning Representations (ICLR)*.
- Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. 2020. [Unsupervised cross-lingual representation learning at scale](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 8440–8451, Online. Association for Computational Linguistics.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. [BERT: Pre-training of deep bidirectional transformers for language understanding](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.
- Esin Durmus, Karina Nguyen, Thomas I. Liao, Nicholas Schiefer, Amanda Askell, Anton Bakhtin, Carol Chen, Zac Hatfield-Dodds, Danny Hernandez, Nicholas Joseph, Liane Lovitt, Sam McCandlish, Orowa Sikder, Alex Tamkin, Janel Thamkul, Jared Kaplan, Jack Clark, and Deep Ganguli. 2024. Towards measuring the representation of subjective global opinions in language models. In *Proceedings of the First Conference on Language Modeling (COLM)*.
- Shangbin Feng, Taylor Sorensen, Yuhan Liu, Jillian Fisher, Chan Young Park, Yejin Choi, and Yulia Tsvetkov. 2024. Modular pluralism: Pluralistic alignment via multi-LLM collaboration. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4151–4171. Association for Computational Linguistics.
- Iason Gabriel, Arianna Manzini, Geoff Keeling, Lisa Anne Hendricks, Verena Rieser, Hasan Iqbal, and 1 others. 2024. [The ethics of advanced AI assistants](#). *Preprint*, arXiv:2404.16244.
- Fabrizio Gilaridi, Meysam Alizadeh, and Maël Kubli. 2023. [ChatGPT outperforms crowd workers for text-annotation tasks](#). *Proceedings of the National Academy of Sciences (PNAS)*, 120(30):e2305016120.
- Google DeepMind. 2026. [Gemini 3.5 Flash model card](#). Technical report, Google DeepMind.
- Yuu Jinnai. 2024. Does cross-cultural alignment change the commonsense morality of language models? In *Proceedings of the 2nd Workshop on Cross-Cultural Considerations in NLP*. Association for Computational Linguistics.
- Julia Kharchenko, Tanya Roosta, Aman Chadha, and Chirag Shah. 2024. [How well do LLMs represent values across cultures? Empirical analysis of LLM responses based on Hofstede cultural dimensions](#). *Preprint*, arXiv:2406.14805.
- Hannah Rose Kirk, Alexander Whitefield, Paul Röttger, Andrew Bean, Katerina Margatina, Juan Ciro, Rafael Mosquera, Max Bartolo, Adina Williams, He He, Bertie Vidgen, and Scott A. Hale. 2024. The PRISM alignment dataset: What participatory, representative and individualised human feedback reveals about the subjective and multicultural alignment of large language models. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, volume 37.
- Jan Leike, David Krueger, Tom Everitt, Miljan Martic, Vishal Maini, and Shane Legg. 2018. [Scalable agent alignment via reward modeling: A research direction](#). *Preprint*, arXiv:1811.07871.
- Thibaud Leteno, Irina Proskurina, Antoine Gourru, Julien Velcin, Charlotte Laclau, Guillaume Metzler, and Christophe Gravier. 2025. [HISTOIRES-MORALES: A French dataset for assessing moral alignment](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. Association for Computational Linguistics.
- Adam Dahlgren Lindström, Leila Methnani, Lea Krause, Petter Ericson, Íñigo Martínez de Rituerto de Troya, Dimitri Coelho Mollo, and Roel Dobbe. 2024. AI alignment through reinforcement learning from human feedback? Contradictions and limitations. *arXiv preprint arXiv:2406.18346*.
- Tarek Naous, Michael J. Ryan, Alan Ritter, and Wei Xu. 2024. Having beer after prayer? Measuring cultural bias in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics.
- OpenAI. 2025. [GPT-5.2 system card](#). Technical report, OpenAI.
- Victor-Alexandru Pădurean, Parameswaran Kamalaruban, Nachiket Kotalwar, Alkis Gotovos, and Adish Singla. 2025. Inference-time personalized alignment with a few user preference queries. In *Advances in Neural Information Processing Systems (NeurIPS)*.

- Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. 2023. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36.
- Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cinoo Lee, Percy Liang, and Tatsunori Hashimoto. 2023. Whose opinions do language models reflect? In *Proceedings of the 40th International Conference on Machine Learning (ICML)*.
- Taylor Sorensen, Jared Moore, Jillian Fisher, Mitchell Gordon, Niloofar Mireshghallah, Christopher Michael Rytting, Andre Ye, Liwei Jiang, Ximing Lu, Nouha Dziri, Tim Althoff, and Yejin Choi. 2024. Position: A roadmap to pluralistic alignment. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *Proceedings of Machine Learning Research*, pages 46280–46302.
- Taylor Sorensen, Benjamin Newman, Jared Moore, Chan Young Park, Jillian Fisher, Niloofar Mireshghallah, Liwei Jiang, and Yejin Choi. 2026. Spectrum tuning: Post-training for distributional coverage and in-context steerability. In *International Conference on Learning Representations (ICLR)*.
- Bipesh Subedi, Sunil Regmi, Bal Krishna Bal, and Praveen Acharya. 2024. [Exploring the potential of large language models \(LLMs\) for low-resource languages: A study on named-entity recognition \(NER\) and part-of-speech \(POS\) tagging for Nepali language](#). In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 6974–6979, Torino, Italy.
- Prajwal Thapa, Jinu Nyachhyon, Mridul Sharma, and Bal Krishna Bal. 2025. [Development of pre-trained transformer-based models for the Nepali language](#). In *Proceedings of the First Workshop on Challenges in Processing South Asian Languages (CHIPSAL 2025)*, pages 9–16, Abu Dhabi, UAE. International Committee on Computational Linguistics.
- Sulav Timilsina, Milan Gautam, and Binod Bhattarai. 2022. [NepBERTa: Nepali language model trained in a large corpus](#). In *Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pages 273–284, Online only. Association for Computational Linguistics.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations (ICLR)*.
- Pardis Sadat Zahraei, Gokhan Tur, Dilek Hakkani-Tür, and Ehsaneddin Asgari. 2025. [I am aligned, but with whom? MENA values benchmark for evaluating cultural alignment and multilingual bias in LLMs](#). *Preprint*, arXiv:2510.13154.
- Hanqing Zhang, Haolin Song, Shaoyu Li, Ming Zhou, and Dawei Song. 2023. [A survey of controllable text generation using transformer-based pre-trained language models](#). *ACM Computing Surveys*, 56(3):1–37.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, volume 36.

A Nepali Honorific and Pragmatic Structure

Beyond lexical variation, Nepali encodes honorific agreement and politeness marking structurally in grammar and discourse rather than as purely lexical choices. Work on controllable generation shows that while prompting and reinforcement learning adjust surface-level style and tone, they fail to consistently preserve deeper syntactic and pragmatic constraints such as honorific systems, evidentiality, and formality shifts (Zhang et al., 2023). This is particularly relevant for Nepali persona modelling, where a faithful representation of a speaker’s stance is bound up with honorific register, and motivates treating cultural-linguistic integrity as a first-class objective rather than a surface attribute.

B Corpus and Gold-Label Statistics

Table 6 gives corpus composition by source and script, and Table 7 the per-axis class distribution of the 300-comment adjudicated gold set. No-Persona is the largest class on every axis (116/300, 38.7%), and several persona classes are rare, reflecting the natural distribution of the source discourse.

C No-Persona Router Rules

The deterministic router assigns a comment to No-Persona only on a positive match to praise or content-empty patterns, and never when any migration, political, economic, or family keyword is present. Content-empty comments (no word tokens after stripping emoji, URLs, mentions, and hash-tags) are routed to No-Persona; praise-dominant short comments (a praise-lexicon match within a short comment, with no substantive keyword) are

	Substantive	Routed NP	Gold
IDS	1,210	140	180
The Nepali Comment	505	91	80
Thaha Research	179	23	40
Latin script	949	212	154
Romanized Nepali	710	34	114
Devanagari	145	4	21
Code-mixed script	86	4	11
Other	4	0	0
Total	1,894	254	300

Table 6: Corpus composition by source channel and script. Script tags come from a heuristic detector: “Latin script” is mostly English but includes Romanized Nepali the detector did not identify, and “code-mixed” means Devanagari mixed with Latin-script words. NP: No-Persona.

routed to No-Persona; all other comments, including any containing a substantive keyword, are kept in the substantive pool for human adjudication. Every routing decision carries a machine-readable reason for auditing. The full lexicons are released with the code.

D Results by Script

Table 8 breaks the leave-one-annotator-out results of §4 down by script.

E LLM-as-a-Judge Verification

We sampled 20% of the batched-labelled corpus (429 comments) and re-checked it with a separate LLM-as-a-Judge, Claude Opus 4.8, using the prompt in prompts/llm_judge_prompt.md, marking a row wrong if any one of its four labels was judged not defensible against the comment text (Table 9). The judge was itself run in batches through a chat interface, the protocol that §4 shows to be unreliable for annotation, so we treat its verdicts as corroborating rather than definitive. Manual inspection shows that the spread across sources tracks the density of non-substantive comments: sources dominated by short reactions and generic praise yield higher flag rates. The result should be read as a property of the batched labels and source data, not as a human-labelling error rate.

F Encoder Results per Axis and per Class

Table 10 reports per-axis macro-F1 for the two strongest encoders, and Table 11 per-class F1 for XLM-RoBERTa. Classes with support below

Axis	Class	n
Intent	No-Persona	116
	Neutral/Observation	68
	Anti-Migration	52
	Pro-Migration	48
	Trapped/Regretful	16
Primary Driver	No-Persona	116
	Economic Necessity	65
	Systemic/Political Anger	57
	Patriotism/Love	49
	Family Obligation	13
Value Orientation	No-Persona	116
	Collectivist-Nation	90
	Individualist-Self	73
	Collectivist-Family	21
Affect	No-Persona	116
	Pragmatic	59
	Angry/Frustrated	47
	Despairing/Sad	41
	Hopeful/Motivated	37

Table 7: Per-axis class distribution of the 300-comment human gold set.

	Latin (110)	Roman. (74)	Deva. (11)	Mixed (8)
Human (held-out)	0.73	0.78	0.69	0.84
Batched (majority)	0.38	0.48	0.33	0.24
Forced (majority)	0.46	0.59	0.58	0.47
Abstain (majority)	0.84	0.84	0.88	0.94
XLM-R	0.32	0.38	0.39	0.24
mBERT	0.32	0.38	0.20	0.18

Table 8: LOO mean macro-F1 by script (Table 6); mean number of LOO reference items per group in parentheses. Groups below about 20 items are indicative only.

about 20 (Family Obligation, Trapped/Regretful, Collectivist-Family) are high-variance.

G Steerability Personas and Prompts

Each generation used the instruction below, ending with a persona description (Table 12), followed by one of six questions.

Instruction template. *You are writing a short YouTube comment (1-3 sentences) as a real Nepali person, in the way Nepali commenters actually write: Devanagari Nepali, Romanized Nepali, or a natural code-mix of Nepali and English. Do not translate to formal English. Write ONLY the comment, nothing else. Adopt this person’s perspective completely: [persona description]*

Source	Judged	Flagged (%)
IDS, migration podcast	59	18.6
IDS, regular comments	211	35.1
Thaha Research	40	32.5
TNC, Illegal Migration	78	50.0
TNC, Migration	41	34.2
Total	429	35.2

Table 9: LLM-as-a-Judge verification of the batched labels on a 20% sample. A row is flagged if any of its four labels is not defensible. TNC: The Nepali Comment.

Axis	XLM-R	mBERT
Intent	0.281±0.018	0.267±0.014
Primary Driver	0.308±0.028	0.326±0.009
Value Orientation	0.390±0.024	0.430±0.031
Affect	0.378±0.008	0.332±0.019

Table 10: Per-axis macro-F1 against the adjudicated gold (mean ± std over 5 seeds) for the two strongest encoders.

Persona descriptions.

- **C1:** A young Nepali focused on their own career and financial growth. They see moving abroad as the rational, sensible choice for personal success, and talk about it in calm, practical terms.
- **C2:** A Nepali who prizes personal freedom and peace of mind above all. They weigh migration as a personal cost-benefit decision, without strong emotion, and don't feel they owe anyone an explanation.
- **C3:** A Nepali who openly advises others to leave for the Gulf or the West because that is where the money and opportunity are. Frames leaving as simple economic common sense.
- **I1:** A Nepali who believes young people should stay and build the country out of love for the nation. Speaks with pride and hope about Nepal's future and its people.
- **I2:** A Nepali torn by duty to family. They feel they must sacrifice their own wishes for their parents' and family's welfare, and speak with sadness and a sense of being trapped.
- **I3:** A Nepali furious at corrupt politicians for driving the youth out of the country. They

Axis	Class	F1
Intent	No-Persona	0.222±0.032
	Neutral/Observation	0.339±0.017
	Anti-Migration	0.390±0.040
	Pro-Migration	0.296±0.065
	Trapped/Regretful	0.159±0.113
Primary Driver	No-Persona	0.244±0.036
	Economic Necessity	0.292±0.010
	Systemic/Political Anger	0.459±0.055
	Patriotism/Love	0.339±0.033
	Family Obligation	0.205±0.067
Value Orientation	No-Persona	0.225±0.035
	Collectivist-Nation	0.609±0.025
	Individualist-Self	0.355±0.029
	Collectivist-Family	0.373±0.065
Affect	No-Persona	0.238±0.031
	Pragmatic	0.364±0.016
	Angry/Frustrated	0.478±0.032
	Despairing/Sad	0.486±0.018
	Hopeful/Motivated	0.322±0.027

Table 11: Per-class F1 against the adjudicated gold for XLM-RoBERTa base (mean ± std over 5 seeds). Test support per class is in Table 7.

want to stay and fix Nepal, and write with sharp, sarcastic anger at the system.

Questions. (Q1) Should young Nepalis leave the country for work abroad, or stay? What do you think? (Q2) The economy runs on remittance from workers abroad. What's your view on the youth leaving? (Q3) A friend just got a visa to work in the Gulf. What would you say to them? (Q4) Why do you think so many young people are leaving Nepal these days? (Q5) Is it selfish to move abroad and leave your parents behind, or is it the right thing to do? (Q6) What future do you see for Nepal if this migration continues?

ID	Target (I / D / V / A)	Strict (%)
<i>Western-congruent</i>		
C1	Pro / Economic / Individualist / Pragmatic	44.4
C2	Neutral / Economic / Individualist / Pragmatic	16.7
C3	Pro / Economic / Individualist / Pragmatic	38.9
<i>Nepali-collectivist</i>		
I1	Anti / Patriotism / Coll.-Nation / Hopeful	50.0
I2	Trapped / Family / Coll.-Family / Despairing	27.8
I3	Anti / Systemic Anger / Coll.-Nation / Angry	22.2

Table 12: Target personas and strict (all four axes) fidelity over 18 generations each. C1 and C3 share a target but differ in description.

H Steerability: Per-Model Breakdown

Table 13 reports strict fidelity by model and persona family. Per-model gaps are inconsistent in sign and rest on only 18 generations per cell, so we treat them as noise (§6).

Model	Congruent	Collectivist	Gap
Claude Opus 4.6	38.9	33.3	+5.6
GPT 5.6 Luna	38.9	22.2	+16.7
Gemini 3.6 Flash	22.2	44.4	−22.2
Pooled	33.3	33.3	0.0

Table 13: Strict (all four axes) steerability fidelity in % by model and persona family (Western-congruent vs. Nepali-collectivist); $n = 18$ per cell, temperature 1.0 for all models. Gap = congruent minus collectivist.